

Bayesian hierarchical models for NHPP using **rstan**

Miao Cai miao.cai@slu.edu

2019-08-08

Contents

1	Model setting	2
2	Simulating data	3
2.1	Theoretical data generating process (DGP)	3
2.2	R code to simulate data and parameters according to the DGP	4
2.3	Generate NHPP data to pass to rstan	5
3	Stan code	6
4	Estimated results	7
4.1	A single simulation to demonstrate	7
4.2	Scale up simulation	8
5	Further improvement	9

1 Model setting

Let $T_{d,s,i}$ denote the time to the d -th driver's s -th shift's i -th critical event. The total number critical events of d -th driver's s -th shift is $n_{d,s}$. The ranges of these notations are:

- $i = 1, 2, \dots, n_{d,S_d}$,
- $s = 1, 2, \dots, S_d$,
- $d = 1, 2, \dots, D$.

We assume the times of critical events within the d -th driver's s -th shift were generated from a non-homogeneous Poisson process (NHPP) with a power law process (PLP), with a fix shape parameter β and varying scale parameters $\theta_{d,s}$ across drivers. The data generating process is then:

$$\begin{aligned}
 T_{d,s,1}, T_{d,s,2}, \dots, T_{d,s,n_{d,s}} &\sim \text{PLP}(\beta, \theta_{d,s}) \\
 \beta &\sim \text{Gamma}(1, 1) \\
 \log \theta_{d,s} &= \gamma_{0d} + \gamma_1 x_{d,s,1} + \gamma_2 x_{d,s,2} + \dots + \gamma_k x_{d,s,k} \\
 \gamma_{01}, \gamma_{02}, \dots, \gamma_{0D} &\sim \text{i.i.d. } N(\mu_0, \sigma_0^2) \\
 \gamma_1, \gamma_2, \dots, \gamma_k &\sim \text{i.i.d. } N(0, 10^2) \\
 \mu_0 &\sim N(0, 10^2) \\
 \sigma_0 &\sim \text{Gamma}(1, 1)
 \end{aligned}$$

2 Simulating data

2.1 Theoretical data generating process (DGP)

1. Random intercepts $\gamma_{01}, \gamma_{02}, \dots, \gamma_{0D}$. The standard deviation of μ_0 was intentionally set to small number 2 to make $\theta_{d,s}$ fall into a reasonably small range. If I otherwise set it as 10, $\theta_{d,s}$ may be more than 10^5 due to the exponentiation, which may not be realistic in real-life data.

$$\begin{aligned}\mu_0 &= 0, \quad \sigma_0 = 0.5 \\ \sigma_0 &\sim \text{Gamma}(1, 1) \\ \gamma_{01}, \gamma_{02}, \dots, \gamma_{0D} &\sim \text{i.i.d. } N(\mu_0, \sigma_0^2)\end{aligned}$$

2. Fixed parameters: 3 fixed parameters $\gamma_1, \gamma_2, \gamma_3$.

$$\gamma_1, \gamma_2, \gamma_3 \sim \text{i.i.d. } N(0, 0.5^2)$$

3. The number of observations in the d -th driver: N_d .

$$N_d \sim \text{Poisson}(10)$$

4. Data: 3 predictor variables $x_{d,s,1}, x_{d,s,2}, x_{d,s,3}$.

$$\begin{aligned}x_{d,s,1} &\sim N(0, 1) \\ x_{d,s,2} &\sim \text{Gamma}(1, 1) \\ x_{d,s,3} &\sim \text{Poisson}(0.2)\end{aligned}$$

5. Scale parameters of a NHPP (random effects): $\theta_{d,s}$.

$$\theta_{d,s} = \text{EXP}(\gamma_{0d} + \gamma_1 x_{d,s,1} + \gamma_2 x_{d,s,2} + \gamma_k x_{d,s,3})$$

6. Shape parameter of a NHPP (fixed effect): $\beta \sim \text{Gamma}(1, 1)$. Set

$$\beta = 1.5$$

7. Simulate truncate time τ_s for each shift.

$$\tau_s \sim N(10, 1.3)$$

8. Simulate a NHPP based on β and $\theta_{d,s}$.

$$T_{d,s,1}, T_{d,s,2}, \dots, T_{d,s,n_{d,s}} \sim \text{PLP}(\beta, \theta_{d,s})$$

2.2 R code to simulate data and parameters according to the DGP

```
pacman::p_load(rstan, tidyverse, data.table)
source("functions/NHPP_functions.R")

set.seed(123)
D = 10 # the number of drivers
K = 3 # the number of predictor variables

# 1. Random-effect intercepts
# hyperparameters
mu0 = 0
sigma0 = 0.5
r_OD = rnorm(D, mean = mu0, sd = sigma0)

# 2. Fixed-effects parameters
R_K = rnorm(K, mean = 0, sd = 0.5)

# 3. The number of observations (shifts) in the $d$-th driver: $N_{\{d\}}$
N_K = rpois(D, 10)
N = sum(N_K) # the total number of obs
id = rep(1:D, N_K)

# 4. Generate data: $x_1, x_2, \dots x_K$
sim1 = function(group_sizes = N_K){
  ntot = sum(group_sizes)

  int1 = rep(1, ntot)
  x1 = rnorm(ntot, 0, 1)
  x2 = rgamma(ntot, 1, 1)
  x3 = rpois(ntot, 0.2)

  return(data.frame(int1, x1, x2, x3))
}
X = sim1(N_K)

# 5. Scale parameters of a NHPP
# 5a. parameter matrix: P
P = cbind(r0 = rep(r_OD, N_K), t(replicate(N, R_K)))
M_logtheta = P*X

# returned parameter for each observed shift
beta = 1.5
theta = exp(rowSums(M_logtheta))
round(theta, 3)

## [1] 0.284 1.721 1.440 0.403 1.672 1.265 0.795 2.269 1.485 1.498 2.208
## [12] 1.502 0.875 0.817 0.713 0.777 1.063 0.626 9.235 4.786 1.793 3.045
## [23] 2.627 3.822 2.729 2.755 2.249 2.171 6.624 2.214 6.020 0.975 3.326
## [34] 1.140 2.079 1.334 1.228 0.852 0.792 0.560 1.887 1.434 1.364 2.426
## [45] 4.139 0.976 0.270 5.343 1.880 2.676 4.558 3.342 1.340 3.739 1.500
## [56] 1.794 1.804 1.680 1.937 1.783 0.728 2.232 0.790 0.474 1.108 1.203
## [67] 0.910 0.646 0.507 1.667 0.597 2.714 1.961 0.772 0.441 0.662 1.144
## [78] 0.740 0.659 0.568 0.916 0.606
```

```
round(r_OD, 3)
```

```
## [1] -0.280 -0.115  0.779  0.035  0.065  0.858  0.230 -0.633 -0.343 -0.223
```

2.3 Generate NHPP data to pass to rstan

```
sim_hier_plp_tau = function(){
  t_list = list()
  len_list = list()
  tau_vector = rnorm(N, 10, 1.3)

  for (i in 1:N) {
    t_list[[i]] = sim_plp_tau(tau_vector[i], beta, theta[i])
    len_list[[i]] = length(t_list[[i]])
  }

  event_dat = data.frame(
    shift_id = rep(1:N, unlist(len_list)),
    event_time = Reduce(c, t_list)
  )

  start_end_dat = data.frame(
    shift_id = 1:N,
    start_time = rep(0, N),
    end_time = tau_vector #difference2
  )

  return(list(event_dat = event_dat,
             start_end_dat = start_end_dat,
             shift_length = unlist(len_list)))
}

df = sim_hier_plp_tau()

hier_dat = list(
  N = nrow(df$event_dat),
  K = nrow(df$start_end_dat),
  D = id, #driver index
  tau = df$start_end_dat$end_time,
  event_time = df$event_dat$event_time,
  s = df$shift_length, #the number of events in each shift
  x1 = X[,2], x2 = X[,3], x3 = X[,4]
)
```

3 Stan code

```
functions{
  real nhpp_log(vector t, real beta, real theta, real tau){
    vector[num_elements(t)] loglik_part;
    real loglikelihood;
    for (i in 1:num_elements(t)){
      loglik_part[i] = log(beta) - beta*log(theta) + (beta - 1)*log(t[i]);
    }
    loglikelihood = sum(loglik_part) - (tau/theta)^beta;
    return loglikelihood;
  }
}

data {
  int<lower=0> N; //total # of obs
  int<lower=0> K; //total # of shifts
  int<lower=0> D[K]; //driver index, this must be an array
  vector<lower=0>[K] tau; //truncated time
  vector<lower=0>[N] event_time; //failure time
  int s[K]; //group sizes
  vector[K] x1;
  vector[K] x2;
  vector[K] x3;
}

parameters{
  real<lower=0> beta;
  vector[K] r0; // random intercept
  vector[3] r; // fixed parameters
  real mu0; // hyperparameter
  real<lower=0> sigma0; // hyperparameter
}

transformed parameters{
  vector<lower=0>[K] theta;
  for (k0 in 1:K){
    theta[k0] = exp(r0[ D[k0] ] + x1[k0]*r[1] + x2[k0]*r[2] + x3[k0]*r[3]);
  }
}

model{
  int position;
  position = 1;
  for (k in 1:K){
    if(s[k] == 0) continue;
    segment(event_time, position, s[k]) ~ nhpp(beta, theta[k], tau[k]);
    position = position + s[k];
  }
  beta ~ gamma(1, 1);
  r0 ~ normal(mu0, sigma0);
  r ~ normal(0, 10);
  mu0 ~ normal(0, 10);
  sigma0 ~ gamma(1, 1);
  theta ~ gamma(1, 0.01);
}
```

4 Estimated results

4.1 A single simulation to demonstrate

```
f = stan("stan/nhpp_plp_tau_ML.stan",
        chains = 1, iter = 1000, data = hier_dat, refresh = 0)

## DIAGNOSTIC(S) FROM PARSER:
## Info:
## Left-hand side of sampling statement (~) may contain a non-linear transform of a parameter or local variable
## If it does, you need to include a target += statement with the log absolute determinant of the Jacobian
## Left-hand-side of sampling statement:
##      theta ~ gamma(...)

## Warning: There were 1 chains where the estimated Bayesian Fraction of Missing Information was low. See
## http://mc-stan.org/misc/warnings.html#bfmi-low

## Warning: Examine the pairs() plot to diagnose sampling problems

## Warning: Bulk Effective Samples Size (ESS) is too low, indicating posterior means and medians may be biased
## Running the chains for more iterations may help. See
## http://mc-stan.org/misc/warnings.html#bulk-ess

## Warning: Tail Effective Samples Size (ESS) is too low, indicating posterior variances and tail quantiles may be biased
## Running the chains for more iterations may help. See
## http://mc-stan.org/misc/warnings.html#tail-ess

pacman::p_load(magrittr)
est = broom::tidy(f)

pull_est = function(var = "theta", est_obj = f){
  z = est_obj %>%
    broom::tidy() %>%
    filter(grepl(var, term)) %>%
    pull(estimate) %>%
    round(3)
  return(z)
}
```

Estimated values:

- Hyperparameters: $\hat{\mu}_0$: 0.043, $\hat{\sigma}_0$: 0.572
- Individual level parameters: $\gamma_1, \gamma_2, \gamma_3$: 0.626, 0.16, 0.164
- Rate parameter β : 1.499
- θ : 0.292, 1.756, 1.416, 0.417, 1.796, 1.308, 0.816, 2.228, 1.567, 1.572, 2.176, 1.529, 0.899, 0.829, 0.73, 0.768, 1.056, 0.605, 9.656, 4.971, 1.69, 2.911, 2.539, 3.928, 2.675, 2.812, 2.295, 2.221, 6.722, 2.225, 6.253, 0.963, 3.417, 1.15, 1.977, 1.351, 1.167, 0.856, 0.747, 0.555, 1.78, 1.449, 1.323, 2.403, 4.42, 0.991, 0.272, 5.425, 1.862, 2.518, 4.72, 3.214, 1.322, 3.689, 1.433, 1.724, 1.788, 1.555, 1.947, 1.66, 0.734, 2.118, 0.796, 0.475, 1.14, 1.195, 0.911, 0.655, 0.481, 1.632, 0.557, 2.695, 1.911, 0.712, 0.446, 0.657, 1.155, 0.753, 0.674, 0.562, 0.926, 0.596

4.2 Scale up simulation

To be added.

5 Further improvement

In Stan code:

- Need a data matrix X ,
- Need matrix multiplication,

In data:

- Need a driver index $d = 1, 2, \dots, K$ for each shift k
- Need a data matrix X

```
pacman::p_load(rstan, tidyverse, data.table)
source("functions/NHPP_functions.R")

sim_hier_nhpp = function(group_size_lambda = 10, D = 10, K = 3, beta = 1.5)
{
  # 1. Random-effect intercepts
  # hyperparameters
  mu0 = 0.2
  sigma0 = 0.5
  r_OD = rnorm(D, mean = mu0, sd = sigma0)

  # 2. Fixed-effects parameters
  R_K = c(1, 0.3, 0.2)

  # 3. The number of shifts in the  $d$ -th driver:  $N_{\{d\}}$ 
  N_K = rpois(D, group_size_lambda)
  N = sum(N_K) # the total number of obs
  id = rep(1:D, N_K)

  # 4. Generate data:  $x_1, x_2, \dots, x_K$ 
  sim1 = function(group_sizes = N_K)
  {
    ntot = sum(group_sizes)
```

```

int1 = rep(1, ntot)
x1 = rnorm(ntot, 1, 1)
x2 = rgamma(ntot, 1, 1)
x3 = rpois(ntot, 2)

return(data.frame(int1, x1, x2, x3))
}
X = sim1(N_K)

# 5. Scale parameters of a NHPP
# 5a. parameter matrix: P
P = cbind(r0 = rep(r_OD, N_K),
          t(replicate(N, R_K)))
M_logtheta = P*X

# returned parameter for each observed shift

theta_vec = exp(rowSums(M_logtheta))

df = sim_hier_plp_tau(N = N, beta = beta, theta = theta_vec)

hier_dat = list(
  N = nrow(df$event_dat),
  K = K,
  S = nrow(df$start_end_dat),
  D = max(id),
  id = id, #driver index
  tau = df$start_end_dat$end_time,
  event_time = df$event_dat$event_time,
  group_size = df$shift_length, #the number of events in each shift
  X_predictors = X[,2:4]
)

```

```

true_params = list(
  mu0 = mu0, sigma0 = sigma0,
  r0 = r_OD, r1_rk = R_K,
  beta = beta,
  theta = theta_vec
)

return(list(hier_dat = hier_dat, true_params = true_params))
}

```

sampling from Stan

```

df = sim_hier_nhpp(D = 10, beta = 1.5)
f = stan("stan/nhpp_plp_tau_ML1.stan",
  chains = 1, iter = 1000, data = df$hier_dat)

```

##

SAMPLING FOR MODEL 'nhpp_plp_tau_ML1' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 0.000307 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 3.07 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 1000 [0%] (Warmup)

Chain 1: Iteration: 100 / 1000 [10%] (Warmup)

Chain 1: Iteration: 200 / 1000 [20%] (Warmup)

Chain 1: Iteration: 300 / 1000 [30%] (Warmup)

Chain 1: Iteration: 400 / 1000 [40%] (Warmup)

Chain 1: Iteration: 500 / 1000 [50%] (Warmup)

Chain 1: Iteration: 501 / 1000 [50%] (Sampling)

Chain 1: Iteration: 600 / 1000 [60%] (Sampling)

```
## Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 2.22766 seconds (Warm-up)
## Chain 1: 1.44747 seconds (Sampling)
## Chain 1: 3.67513 seconds (Total)
## Chain 1:
```

```
# check estimation results
f
```

```
## Inference for Stan model: nhpp_plp_tau_ML1.
## 1 chains, each with iter=1000; warmup=500; thin=1;
## post-warmup draws per chain=500, total post-warmup draws=500.
##
```

##	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff
## mu0	1.77	0.01	0.13	1.52	1.69	1.78	1.86	2.02	296
## sigma0	0.35	0.01	0.11	0.21	0.28	0.33	0.39	0.58	320
## beta	1.46	0.00	0.05	1.37	1.42	1.46	1.49	1.56	259
## R1_K[1]	0.92	0.00	0.05	0.82	0.88	0.92	0.96	1.03	195
## R1_K[2]	0.25	0.00	0.05	0.15	0.22	0.25	0.28	0.34	554
## R1_K[3]	0.22	0.00	0.04	0.15	0.19	0.22	0.24	0.30	367
## R0[1]	1.52	0.01	0.10	1.31	1.46	1.52	1.58	1.73	205
## R0[2]	1.74	0.01	0.11	1.52	1.66	1.74	1.81	1.95	182
## R0[3]	1.52	0.01	0.09	1.36	1.46	1.52	1.58	1.70	265
## R0[4]	1.81	0.01	0.11	1.59	1.73	1.81	1.88	2.04	419
## R0[5]	1.65	0.01	0.11	1.44	1.59	1.65	1.72	1.87	202
## R0[6]	1.25	0.01	0.11	1.02	1.17	1.25	1.33	1.47	126
## R0[7]	2.03	0.01	0.17	1.68	1.93	2.02	2.14	2.36	292
## R0[8]	1.92	0.01	0.12	1.70	1.85	1.93	1.99	2.17	460
## R0[9]	2.09	0.01	0.11	1.89	2.02	2.08	2.15	2.31	393
## R0[10]	2.20	0.01	0.13	1.94	2.11	2.19	2.28	2.44	229

## mu0_true	0.19	0.01	0.15	-0.09	0.09	0.19	0.30	0.46	573
## R0_true[1]	-0.06	0.01	0.11	-0.27	-0.14	-0.05	0.02	0.14	473
## R0_true[2]	0.16	0.01	0.12	-0.09	0.08	0.15	0.24	0.37	381
## R0_true[3]	-0.06	0.01	0.12	-0.29	-0.14	-0.06	0.03	0.15	444
## R0_true[4]	0.23	0.01	0.14	-0.04	0.13	0.23	0.32	0.48	433
## R0_true[5]	0.07	0.01	0.12	-0.16	0.00	0.07	0.15	0.31	514
## R0_true[6]	-0.33	0.01	0.12	-0.59	-0.41	-0.33	-0.25	-0.10	348
## R0_true[7]	0.45	0.01	0.20	0.05	0.33	0.45	0.59	0.83	342
## R0_true[8]	0.34	0.01	0.14	0.08	0.25	0.34	0.44	0.61	397
## R0_true[9]	0.51	0.01	0.14	0.21	0.42	0.50	0.60	0.78	598
## R0_true[10]	0.61	0.01	0.14	0.34	0.53	0.62	0.71	0.87	404
## lp_	298.05	0.22	3.04	291.48	296.26	298.25	300.24	303.28	193
##	Rhat								
## mu0	1.00								
## sigma0	1.00								
## beta	1.00								
## R1_K[1]	1.00								
## R1_K[2]	1.00								
## R1_K[3]	1.00								
## R0[1]	1.00								
## R0[2]	1.00								
## R0[3]	1.00								
## R0[4]	1.00								
## R0[5]	1.00								
## R0[6]	1.00								
## R0[7]	1.00								
## R0[8]	1.00								
## R0[9]	1.00								
## R0[10]	1.00								
## mu0_true	1.00								
## R0_true[1]	1.00								
## R0_true[2]	1.01								
## R0_true[3]	1.00								

```

## R0_true[4]  1.00
## R0_true[5]  1.00
## R0_true[6]  1.00
## R0_true[7]  1.00
## R0_true[8]  1.00
## R0_true[9]  1.00
## R0_true[10] 1.00
## lp__        1.00
##
## Samples were drawn using NUTS(diag_e) at Thu Aug  8 00:06:58 2019.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

```

```
df$true_params
```

```

## $mu0
## [1] 0.2
##
## $sigma0
## [1] 0.5
##
## $r0
## [1] 0.007600664 0.365340280 -0.077310225 0.260786157 0.176201942
## [6] -0.188125795 0.615720626 0.623153919 0.712069753 0.833998293
##
## $r1_rk
## [1] 1.0 0.3 0.2
##
## $beta
## [1] 1.5
##
## $theta
## [1] 0.6034702 2.4859378 1.1791556 11.9240479 8.0888078

```

##	[6]	8.2844389	5.9921438	4.0092433	4.0793214	2.4560764
##	[11]	0.4900528	9.1472993	7.1127270	385.3064549	46.8419284
##	[16]	9.5229321	1.4952724	1.3077455	7.0098578	1.2306090
##	[21]	11.2896253	4.3779013	1.9697259	1.7257327	4.6076529
##	[26]	2.9979074	5.0998086	5.4272624	28.3335796	6.9349328
##	[31]	4.4659386	4.2114884	2.3906267	27.5547284	6.2070766
##	[36]	6.4071452	15.5636042	5.6499430	3.5256133	70.0814971
##	[41]	5.5929231	2.7913710	24.2494362	0.9141287	5.0234923
##	[46]	23.1163245	7.8887152	8.6089758	2.1110206	23.9357796
##	[51]	49.6033563	0.3241403	14.0838740	8.6657893	8.7143414
##	[56]	8.5086261	6.7382076	15.6917833	87.6953307	15.6815778
##	[61]	2.5723906	13.2830855	9.8674529	235.5315816	3.3269894
##	[66]	7.3713720	10.7808832	17.6446845	4.2852413	24.0938685
##	[71]	50.1672660	3.4503360	17.0192166	12.2051503	14.3430527
##	[76]	51.6852054	2.9420156	11.3943026	2.6243754	4.3414671
##	[81]	14.2609634	1.2562075	35.9411413	3.6466677	16.6637638
##	[86]	2.4246154	10.4860473	10.7440957	4.9308350	28.8213547
##	[91]	32.5376764	10.0808709	1.3616594	10.4626472	1.6026213
##	[96]	78.7852806	27.0651914	31.7142262	16.8507354	15.4378015
##	[101]	11.4646960				